

Paper: Attention Is All You Need (sample notes)

The paper introduces the Transformer architecture based on self-attention.

It replaces recurrence with multi-head attention and position-wise feed-forward blocks.

Key benefits include parallel training, strong sequence modeling, and improved translation quality.

This sample file is included for quick local RAG testing.